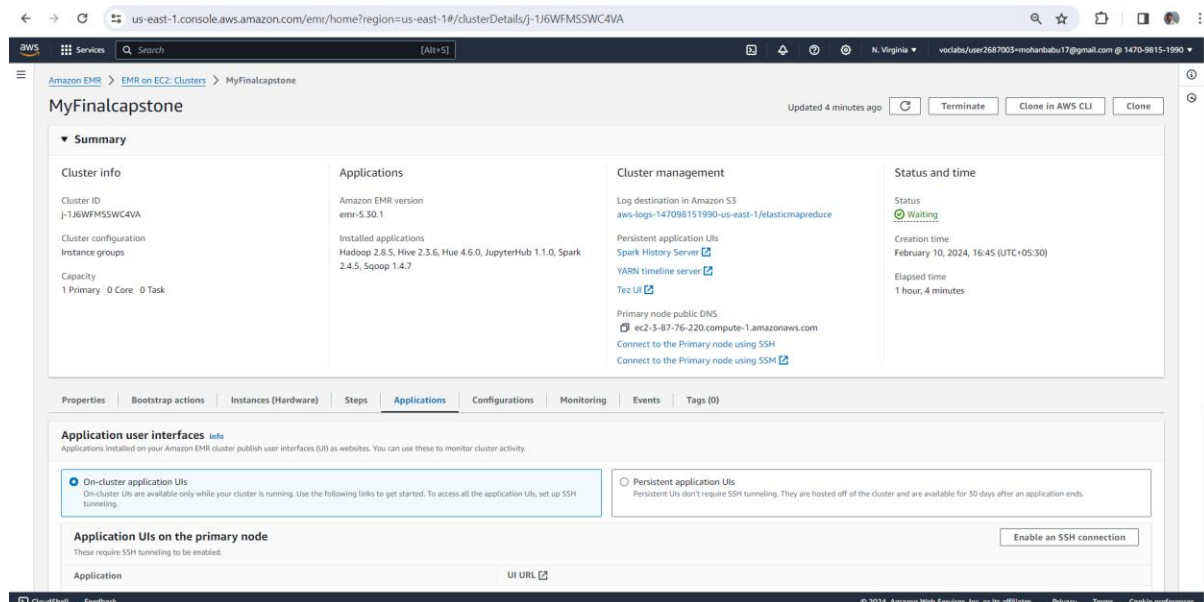


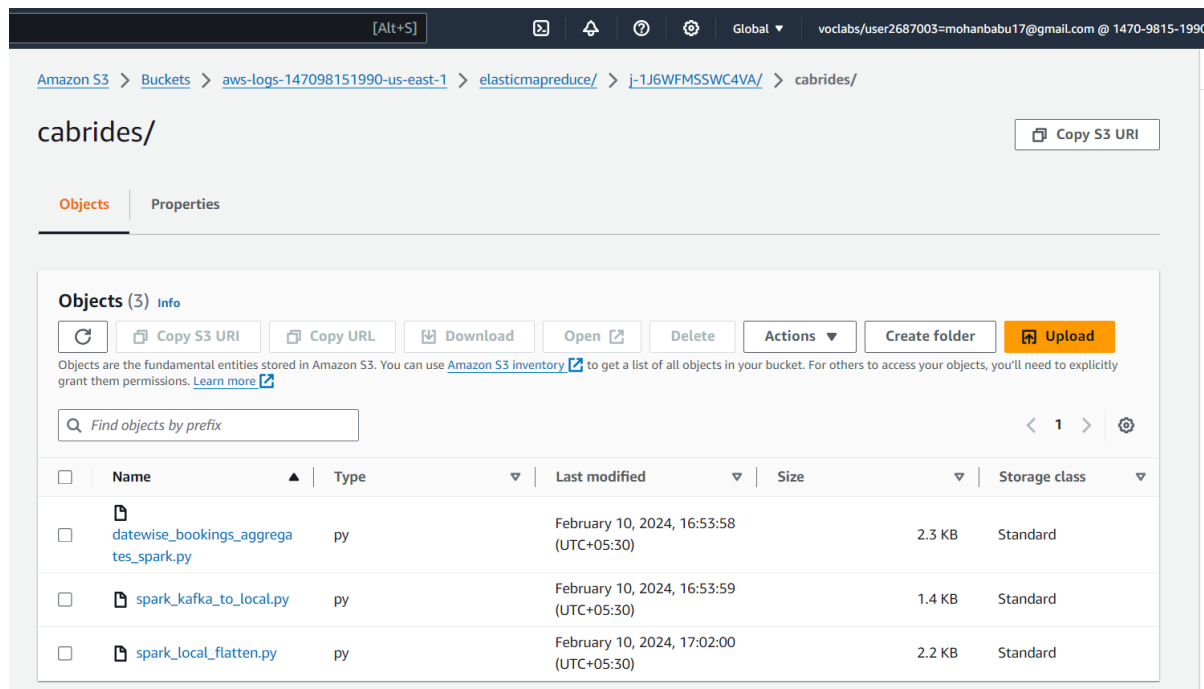
Logic For Final Submission

Active EMR cluster:



The screenshot shows the Amazon EMR console for a cluster named 'MyFinalcapstone'. The cluster is in the 'Waiting' status. The console displays various details including cluster ID, configuration, capacity, and applications installed. The 'Applications' section lists Amazon EMR version, installed applications (Hadoop, Hive, Hue, JupyterHub, Spark, Tez, Sqoop), and cluster management options like log destination, persistent application UIs, and primary node public DNS.

S3 files:



The screenshot shows the Amazon S3 console for a bucket named 'cabrides/'. The console displays a list of objects (files) stored in the bucket. The objects are listed with their names, types, last modified dates, sizes, and storage classes.

Name	Type	Last modified	Size	Storage class
datewise_bookings_aggregates_spark.py	py	February 10, 2024, 16:53:58 (UTC+05:30)	2.3 KB	Standard
spark_kafka_to_local.py	py	February 10, 2024, 16:53:59 (UTC+05:30)	1.4 KB	Standard
spark_local_flatten.py	py	February 10, 2024, 17:02:00 (UTC+05:30)	2.2 KB	Standard

Validation already completed for task 1 to task 4 in LogicFirstSubmission.pdf

Capstone Project - Validation Document

Note: Please note that records retrieved from the Kafka topic and related metrics given below can vary a bit.

Data Ingestion with Sqoop

Please check the number of records that are imported after the Sqoop Job

```
Number of records retrieved - 1000
```

Bookings Table Count

Please check the number of records in the bookings table

```
Number of records - 1000
```

Clickstream Table Count

Please check the number of records in the clickstream table

```
Number of records - 2984
```

Bookings Aggregates Table Count

Please check the number of records in the bookings aggregates table

```
Number of records - 289
```

Enable HUE UI for Hive execution: add port 8888 in both master and slave security group

EC2 > Security Groups > sg-05bc9cf1710692fbd - ElasticMapReduce-master

sg-05bc9cf1710692fbd - ElasticMapReduce-master

Details

Security group name ElasticMapReduce-master	Security group ID sg-05bc9cf1710692fbd	Description Master group for Elastic MapReduce created on 2024-01-06T23:44:34.662Z	VPC ID vpc-0f36e4dc26e9f710a
Owner 147098151990	Inbound rules count 11 Permission entries	Outbound rules count 1 Permission entry	

Inbound rules (11)

Name	Security group rule...	IP version	Type	Protocol	Port range	Source	Description
-	sgr-01c15a096a843ab...	IPv4	SSH	TCP	22	0.0.0.0/0	-
-	sgr-02ae51e8228e1f3c3	-	All TCP	TCP	0 - 65535	sg-05bc9cf1710692fb...	-
-	sgr-0421dc0125cb15fe	IPv4	Custom TCP	TCP	8888	0.0.0.0/0	-
-	sgr-07155bb117763f0...	-	All ICMP - IPv4	ICMP	All	sg-05bc9cf1710692fb...	-
-	sgr-0cee5b818392107...	-	All UDP	UDP	0 - 65535	sg-0ac22a1e3641eea3...	-

EC2 > Security Groups > sg-0ac22a1e3641eea35 - ElasticMapReduce-slave

sg-0ac22a1e3641eea35 - ElasticMapReduce-slave

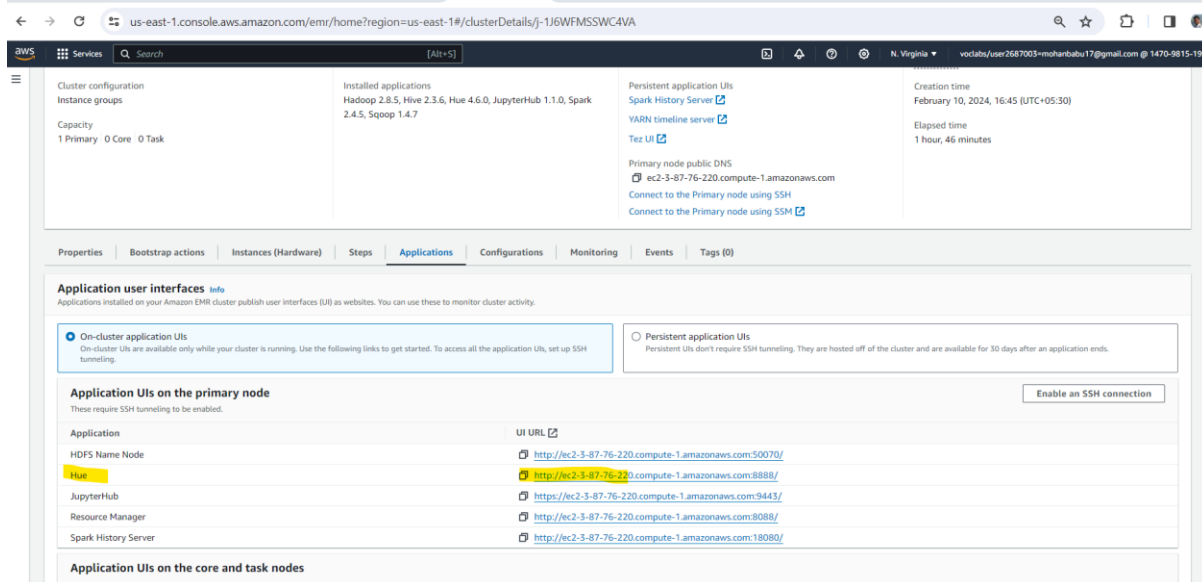
Details

Security group name ElasticMapReduce-slave	Security group ID sg-0ac22a1e3641eea35	Description Slave group for Elastic MapReduce created on 2024-01-06T23:44:34.662Z	VPC ID vpc-0f36e4dc26e9f710a
Owner 147098151990	Inbound rules count 8 Permission entries	Outbound rules count 1 Permission entry	

Inbound rules (8)

Name	Security group rule...	IP version	Type	Protocol	Port range	Source	Description
-	sgr-0d9637f9d68c22bd6	-	All ICMP - IPv4	ICMP	All	sg-0ac22a1e3641eea3...	-
-	sgr-092ee6d556a27b5...	-	All UDP	UDP	0 - 65535	sg-05bc9cf1710692fb...	-
-	sgr-0862616dd8bb49...	IPv4	SSH	TCP	22	0.0.0.0/0	-
-	sgr-087124a5a1b218...	IPv4	Custom TCP	TCP	8888	0.0.0.0/0	-
-	sgr-04288a139d6ea7b...	-	All TCP	TCP	0 - 65535	sg-0ac22a1e3641eea3...	-

Open HUE link:



us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-136WFMSWC4VA

Cluster configuration
Instance groups
Capacity
1 Primary 0 Core 0 Task

Installed applications
Hadoop 2.8.5, Hive 2.3.6, Hue 4.6.0, JupyterHub 1.1.0, Spark 2.4.5, Sqoop 1.4.7

Persistent application UIs
Spark History Server
YARN timeline server
Tez UI
Primary node public DNS
ec2-3-87-76-220.compute-1.amazonaws.com
Connect to the Primary node using SSH
Connect to the Primary node using SSM

Creation time
February 10, 2024, 16:45 (UTC+05:30)
Elapsed time
1 hour, 46 minutes

Properties | Bootstrap actions | Instances (Hardware) | Steps | **Applications** | Configurations | Monitoring | Events | Tags (0)

Application user interfaces
Applications installed on your Amazon EMR cluster publish user interfaces (UI) as webpages. You can use these to monitor cluster activity.

☒ On-cluster application UIs
On-cluster UIs are available only while your cluster is running. Use the following links to get started. To access all the application UIs, set up SSH tunneling.

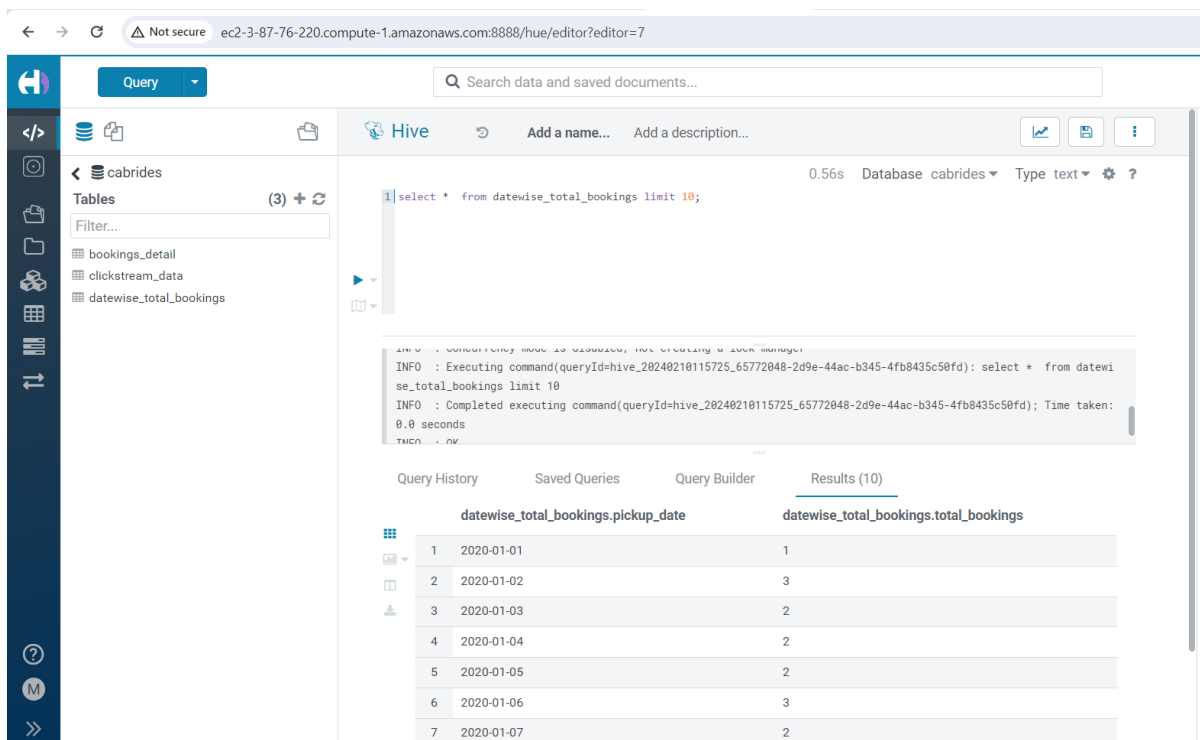
☐ Persistent application UIs
Persistent UIs don't require SSH tunneling. They are hosted off of the cluster and are available for 30 days after an application ends.

Application UIs on the primary node
These require SSH tunneling to be enabled. [Enable an SSH connection](#)

Application	UI URL
HDFS Name Node	http://ec2-3-87-76-220.compute-1.amazonaws.com:50070/
Hue	http://ec2-3-87-76-220.compute-1.amazonaws.com:8888/
JupyterHub	https://ec2-3-87-76-220.compute-1.amazonaws.com:9443/
Resource Manager	http://ec2-3-87-76-220.compute-1.amazonaws.com:8088/
Spark History Server	http://ec2-3-87-76-220.compute-1.amazonaws.com:18080/

Application UIs on the core and task nodes

HUE: query check in UI, select cabrides database and execute table query



Not secure ec2-3-87-76-220.compute-1.amazonaws.com:8888/hue/editor?editor=7

Query

Search data and saved documents...

Hive Add a name... Add a description...

0.56s Database cabrides Type text ?

1 select * from datewise_total_bookings limit 10;

INFO : Executing command(queryId=hive_20240210115725_65772048-2d9e-44ac-b345-4fb8435c50fd): select * from datewise_total_bookings limit 10
INFO : Completed executing command(queryId=hive_20240210115725_65772048-2d9e-44ac-b345-4fb8435c50fd); Time taken: 0.0 seconds

Query History Saved Queries Query Builder **Results (10)**

	datewise_total_bookings.pickup_date	datewise_total_bookings.total_bookings
1	2020-01-01	1
2	2020-01-02	3
3	2020-01-03	2
4	2020-01-04	2
5	2020-01-05	2
6	2020-01-06	3
7	2020-01-07	2

Task 5: Calculate the total number of different drivers for each customer

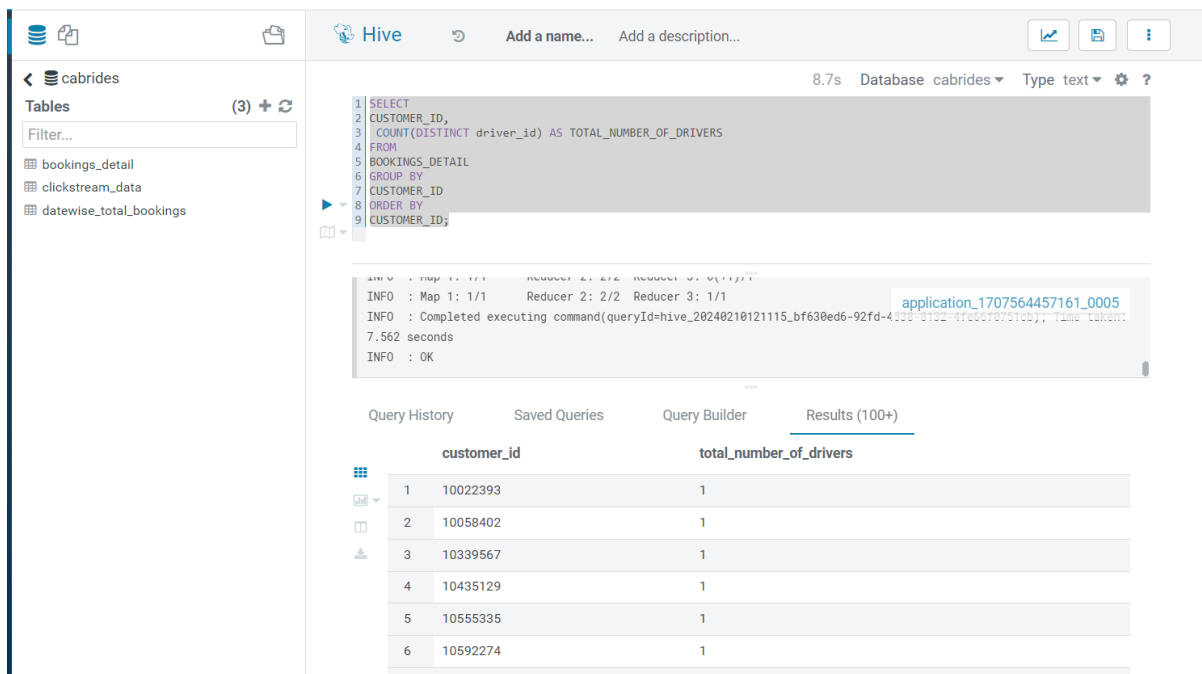
QUERY:

```
SELECT CUSTOMER_ID, COUNT(DISTINCT driver_id) AS TOTAL_NUMBER_OF_DRIVERS
FROM BOOKINGS_DETAIL
GROUP BY CUSTOMER_ID
ORDER BY CUSTOMER_ID;
```

EXPLANATION:

Check multiple drivers for each customer by grouping customer and count total individual drivers for each and sort by customer id

OUTPUT: each customer is having only one driver , so total number of drivers is “1”



The screenshot shows the Hive query interface. The query is executed successfully, and the results are displayed in a table. The table has two columns: 'customer_id' and 'total_number_of_drivers'. The results show that each customer has exactly one driver.

customer_id	total_number_of_drivers
10022393	1
10058402	1
10339567	1
10435129	1
10555335	1
10592274	1

VALIDATION: Exact Match

1. When you run the query to calculate the total number of different drivers for each customer, you would get an output as shown below:

```
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2020-11-17 12:23:06,034 Stage-1 map = 0%, reduce = 0%
2020-11-17 12:23:12,394 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.27 sec
2020-11-17 12:23:20,727 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.69 sec
MapReduce Total cumulative CPU time: 7 seconds 690 msec
Ended Job = job_1605615116654_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 7.69 sec HDFS Read: 43007 HDFS Write: 11000 SUCCESS
Total MapReduce CPU Time Spent: 7 seconds 690 msec
OK
10022393      1
10058402      1
10339567      1
10435129      1
10555335      1
10592274      1
10614890      1
10678994      1
11264797      1
11353346      1
11418437      1
11438890      1
11454977      1
11479815      1
11518953      1
11590321      1
11596512      1
11608791      1
11655671      1
11757536      1
11764909      1
11860278      1
11981042      1
12106105      1
12142182      1
12312603      1
12334699      1
12367832      1
12856708      1
12885363      1
12913608      1
12914577      1
12966909      1
13015449      1
13229062      1
```

• Task 6:

Calculate the total rides taken by each customer.

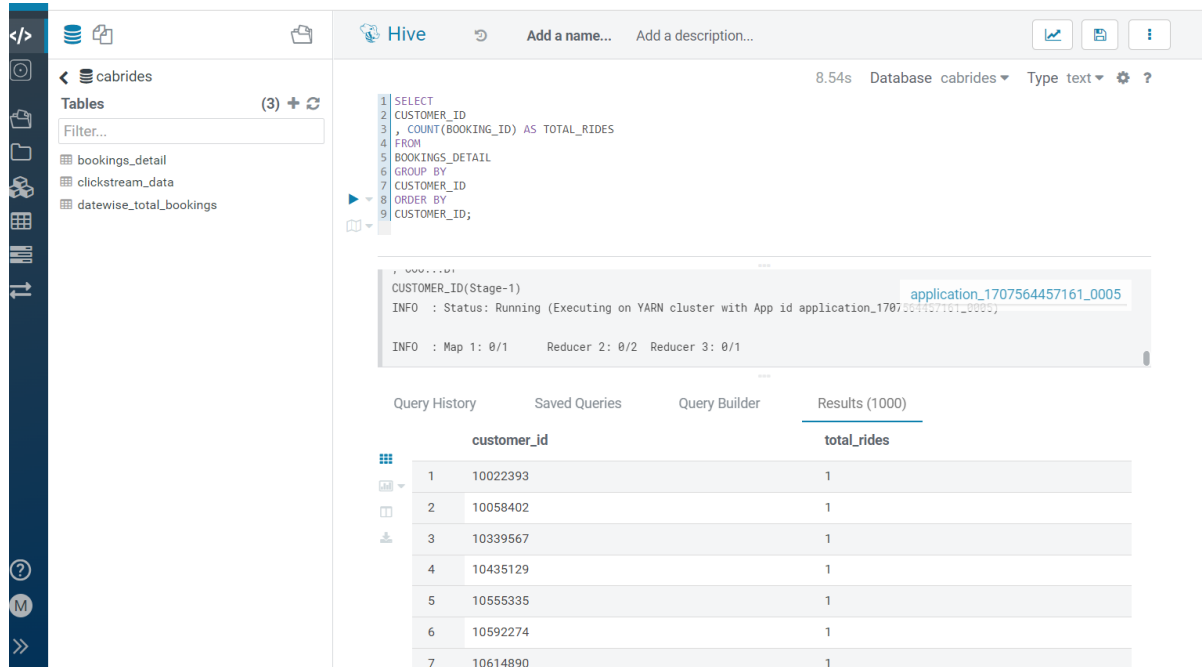
QUERY:

```
SELECT CUSTOMER_ID, COUNT(BOOKING_ID) AS TOTAL_RIDES
FROM BOOKINGS_DETAIL
GROUP BY CUSTOMER_ID
ORDER BY CUSTOMER_ID;
```

EXPLANATION:

Check total rides for each customer by grouping customer, count total rides and sort by customer id

OUTPUT: each customer is having only one ride , so total ride is “1” for each customer



The screenshot shows the Hive query interface. On the left, there's a sidebar with a 'Tables' section showing a list of tables: bookings_detail, clickstream_data, and datewise_total_bookings. The main area displays a SQL query:

```
1 SELECT
2 CUSTOMER_ID
3 , COUNT(BOOKING_ID) AS TOTAL_RIDES
4 FROM
5 BOOKINGS_DETAIL
6 GROUP BY
7 CUSTOMER_ID
8 ORDER BY
9 CUSTOMER_ID;
```

Below the query, there's a status bar indicating the query is running on a YARN cluster with App id application_1707564457161_0005. The results are displayed in a table with two columns: customer_id and total_rides. The results show 7 rows, each with a customer_id and a total_rides value of 1.

customer_id	total_rides
10022393	1
10058402	1
10339567	1
10435129	1
10555335	1
10592274	1
10614890	1

VALIDATION: Exact Match

```
Ended Job = job_1605615116654_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.65 sec HDFS Read: 38721 HDFS Write: 11000 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 650 msec
OK
10022393 1
10058402 1
10339567 1
10435129 1
10555335 1
10592274 1
10614890 1
10678994 1
11264797 1
11353346 1
11418437 1
11438890 1
11454977 1
11479815 1
11518953 1
11580321 1
11596512 1
11608791 1
11655671 1
11757536 1
11764909 1
11860278 1
11981042 1
12106105 1
12142182 1
12312603 1
12331666 1
```

- **Task 7:**

Find the total visits made by each customer on the booking page and the total 'Book Now' button presses. This can show the conversion ratio.

The booking page id is 'e7bc5fb2-1231-11eb-adc1-0242ac120002'.

The Book Now button id is 'fcba68aa-1231-11eb-adc1-0242ac120002'. You also need to calculate the conversion ratio as part of this task. Conversion ratio can be calculated as Total 'Book Now' Button Press/Total Visits made by customer on the booking page.

QUERY:

SELECT

SUM(CASE WHEN PAGE_ID = 'e7bc5fb2-1231-11eb-adc1-0242ac120002' THEN 1 ELSE 0
END) AS

TOTAL_PAGE_VISITS,

SUM(CASE WHEN BUTTON_ID = 'fcba68aa-1231-11eb-adc1-0242ac120002' THEN 1 ELSE 0
END) AS

TOTAL_BUTTON_PRESSED,

ROUND(CAST(SUM(CASE WHEN BUTTON_ID = 'fcba68aa-1231-11eb-adc1-0242ac120002'
THEN 1 ELSE 0 END) AS FLOAT) /

CAST(SUM(CASE WHEN PAGE_ID = 'e7bc5fb2-1231-11eb-adc1-0242ac120002' THEN 1
ELSE

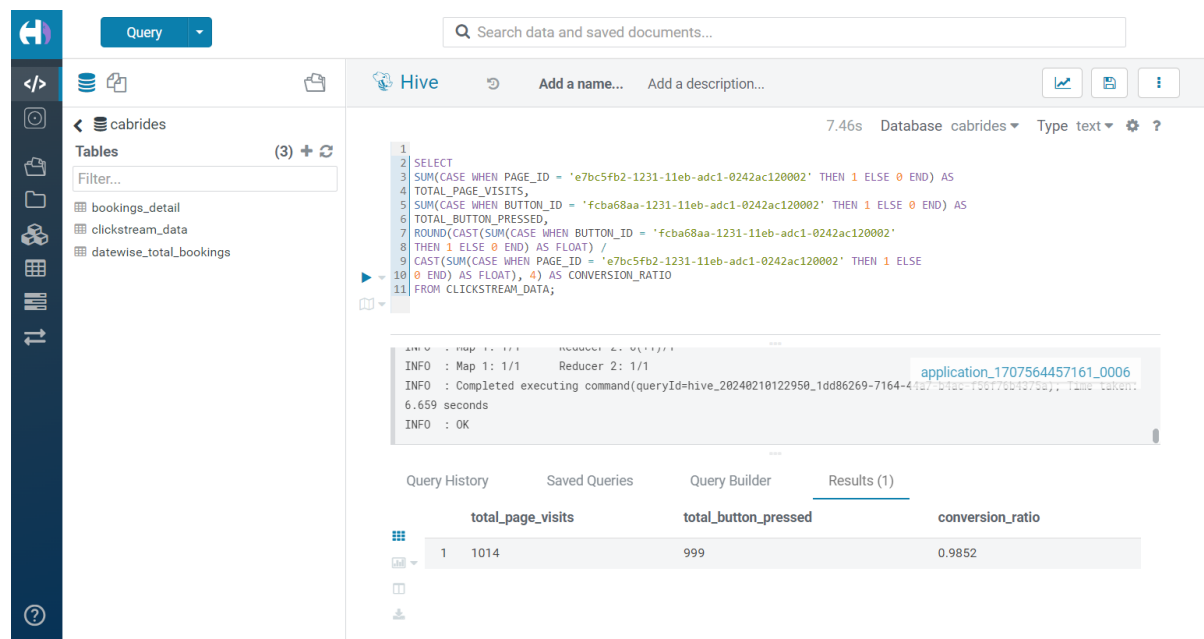
0 END) AS FLOAT), 4) AS CONVERSION_RATIO

FROM CLICKSTREAM_DATA;

EXPLANATION:

- This analysis is useful for understanding customer behavior, such as how often they book a ride after visiting the booking page.
- Tracking the number of times the booking page is visited (PAGE_ID = 'e7bc5fb2-1231-11eb-adc1-0242ac120002') provides the total count of page visits.
- Recording the number of times the booking button is clicked (BUTTON_ID = 'fc6ba68a-1231-11eb-adc1-0242ac120002') indicates the total number of rides that have been booked.
- The conversion ratio is simply the number of total bookings divided by the number of total page visits.
- The conversion ratio is an essential Key Performance Indicator (KPI) for the company, showing a high likelihood of booking when a customer visits the booking page, which is approximately 98% in this instance.

OUTPUT:



The screenshot shows the Hive Query Editor interface. On the left, there's a sidebar with a 'Query' dropdown and a list of tables under the 'cabrides' database: 'bookings_detail', 'clickstream_data', and 'datewise_total_bookings'. The main area displays a SQL query:

```
1 SELECT
2 SUM(CASE WHEN PAGE_ID = 'e7bc5fb2-1231-11eb-adc1-0242ac120002' THEN 1 ELSE 0 END) AS
3 TOTAL_PAGE_VISITS,
4 SUM(CASE WHEN BUTTON_ID = 'fcb68aa-1231-11eb-adc1-0242ac120002' THEN 1 ELSE 0 END) AS
5 TOTAL_BUTTON_PRESSED,
6 ROUND(CAST(SUM(CASE WHEN BUTTON_ID = 'fcb68aa-1231-11eb-adc1-0242ac120002'
7 THEN 1 ELSE 0 END) AS FLOAT) /
8 CAST(SUM(CASE WHEN PAGE_ID = 'e7bc5fb2-1231-11eb-adc1-0242ac120002' THEN 1 ELSE
9 0 END) AS FLOAT), 4) AS CONVERSION_RATIO
10 FROM CLICKSTREAM_DATA;
```

Below the query, the execution status is shown: 'INFO : Map 1: 1/1 Reducer 2: 1/1', 'INFO : Completed executing command(queryId=hive_20240210122950_1dd86269-7164-4027-b100-706770b4370a)', '6.659 seconds', and 'INFO : OK'. The results are displayed in a table with the following data:

total_page_visits	total_button_pressed	conversion_ratio
1014	999	0.9852

VALIDATION: close to Match (0.9852), since Kafka had extra 16 records compare to validation it should get the conversion ratio as 0.9688.

- **Task 8:** Calculate the count of all trips done on black cabs.

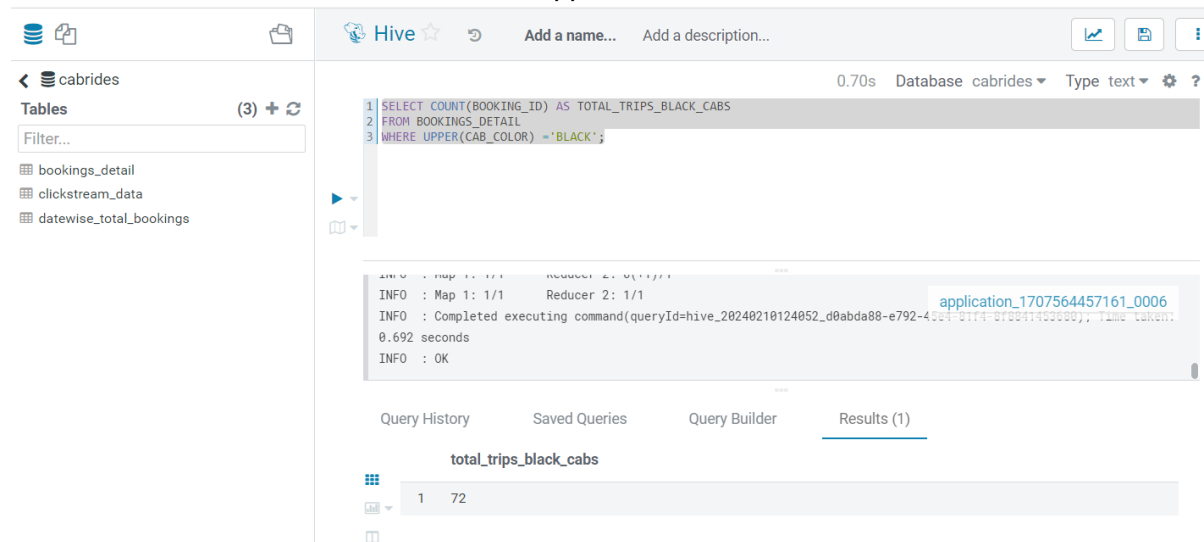
QUERY:

```
SELECT COUNT(BOOKING_ID) AS TOTAL_TRIPS_BLACK_CABS
FROM BOOKINGS_DETAIL
WHERE UPPER(CAB_COLOR) = 'BLACK';
```

EXPLANATION:

- This assessment aids in determining the overall number of journeys made using Black taxis.
- Categorizing by CAB_COLOR can reveal the total count or proportion of trips completed by taxis of a particular color.
- This might reveal any obscure trends in booking related to the color of the taxis, such as a customer's preference for booking taxis of a certain color.

OUTPUT: since case sensitive, handles upper from cab colour



The screenshot shows the Hive query interface. On the left, a sidebar lists tables: bookings_detail, clickstream_data, and datewise_total_bookings. The main area displays a SQL query:

```
1 SELECT COUNT(BOOKING_ID) AS TOTAL_TRIPS_BLACK_CABS
2 FROM BOOKINGS_DETAIL
3 WHERE UPPER(CAB_COLOR) = 'BLACK';
```

Below the query, the execution status is shown: "INFO : Map 1: 1/1 Reducer 2: 1/1", "INFO : Completed executing command(queryId=hive_20240210124052_d0abda88-e792-40e4-b1f4-070041400000), time taken: 0.692 seconds", and "INFO : OK". At the bottom, the results are displayed in a table:

total_trips_black_cabs
72

VALIDATION: Exact Match

Count of all trips done on black cabs -72.

- **Task 9:** Calculate the total amount of tips given date wise to all drivers by customers.

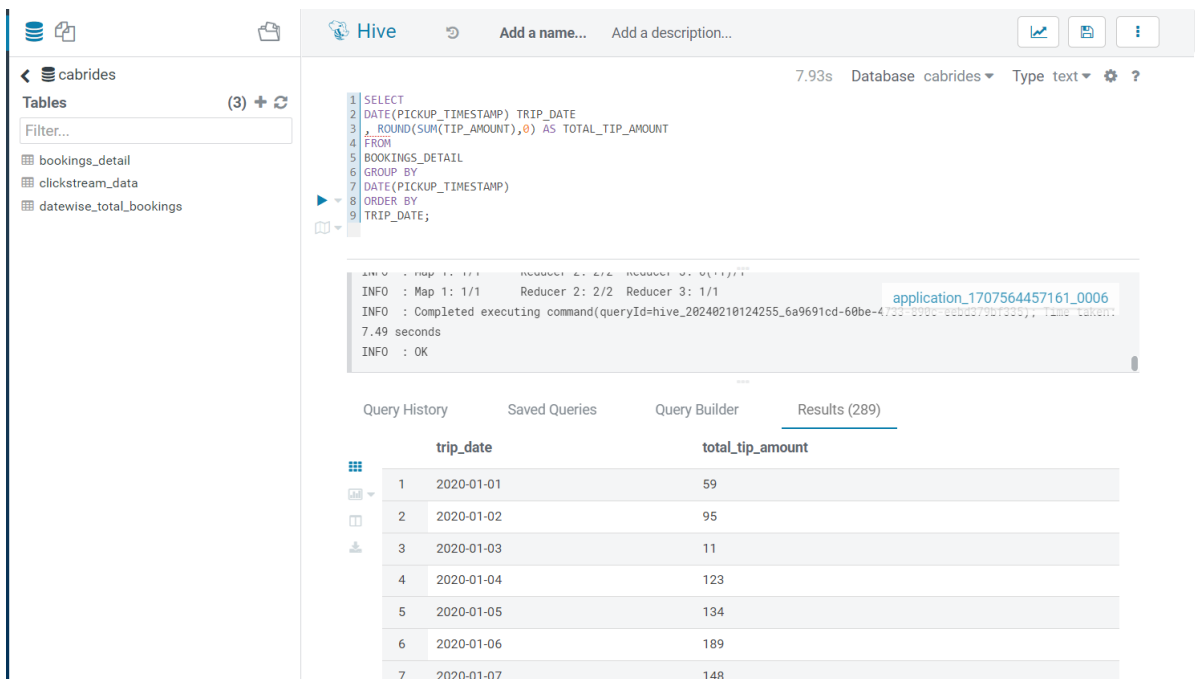
QUERY:

```
SELECT DATE(PICKUP_TIMESTAMP) TRIP_DATE
, ROUND(SUM(TIP_AMOUNT),0) AS TOTAL_TIP_AMOUNT
FROM BOOKINGS_DETAIL
GROUP BY DATE(PICKUP_TIMESTAMP)
ORDER BY TRIP_DATE;
```

EXPLANATION:

- The date function isolates the date from the datetime_stamp value and labels it as Pickup_Date.
- The SUM function, in conjunction with GROUP BY Pickup Date, calculates the aggregate tip amount for each Pickup date.
- The ORDER BY clause, applied to the Pickup Date alias, arranges the results in chronological order based on the date of pickup.
- Analyzing this data can provide insights into whether customers tend to tip more on certain occasions or specific days.
- Based on these insights, management can introduce promotions for customers or drivers, inferring that a higher tip amount correlates with greater customer satisfaction on particular days.

OUTPUT:



The screenshot shows the Hive query interface. On the left, a sidebar lists tables: bookings_detail, clickstream_data, and datewise_total_bookings. The main area displays a SQL query:

```
1 SELECT
2 DATE(PICKUP_TIMESTAMP) TRIP_DATE
3 , ROUND(SUM(TIP_AMOUNT),0) AS TOTAL_TIP_AMOUNT
4 FROM
5 BOOKINGS_DETAIL
6 GROUP BY
7 DATE(PICKUP_TIMESTAMP)
8 ORDER BY
9 TRIP_DATE;
```

Below the query, the execution progress is shown: Map 1: 1/1, Reducer 2: 2/2, Reducer 3: 1/1. The status is "Completed executing command(queryId=hive_20240210124255_6a9691cd-60be-4793-b00c-b0b0370b7039), Time taken: 7.49 seconds". The application ID is application_1707564457161_0006.

At the bottom, the "Results (289)" tab is selected, showing a table with two columns: trip_date and total_tip_amount. The results are as follows:

trip_date	total_tip_amount
2020-01-01	59
2020-01-02	95
2020-01-03	11
2020-01-04	123
2020-01-05	134
2020-01-06	189
2020-01-07	148

VALIDATION: Exact Match

2020-01-01	59
2020-01-02	95
2020-01-03	11
2020-01-04	123
2020-01-05	134
2020-01-06	189
2020-01-07	148
2020-01-08	111
2020-01-09	48
2020-01-10	77
2020-01-11	81
2020-01-12	109
2020-01-14	142
2020-01-15	338
2020-01-16	155
2020-01-17	296
2020-01-18	240
2020-01-20	210
2020-01-21	5
2020-01-23	148
2020-01-24	472
2020-01-25	98
2020-01-26	209
2020-01-27	231
2020-01-28	567

- **Task 10:**

Calculate the total count of all the bookings with ratings lower than 2 as given by customers in a particular month.



QUERY:


```
SELECT DATE_FORMAT(PICKUP_TIMESTAMP, 'yyyy-MM') TRIP_MONTH  
, COUNT(BOOKING_ID) AS NO_OF_BOOKINGS  
FROM BOOKINGS_DETAIL  
WHERE RATING_BY_CUSTOMER < 2  
GROUP BY DATE_FORMAT(PICKUP_TIMESTAMP, 'yyyy-MM')  
ORDER BY TRIP_MONTH;
```

EXPLANATION:

- DATE_FORMAT function formats datetimestamp value in the specified format like yyyy-MM in this case which results like 2023-06.
- WHERE clause is used to filter bookings where rating given by customers is less than 2 which indicates customers dissatisfaction.
- ORDER BY clause with Trip month alias is used to show output in ascending order of pickup month.
- This analysis could help to understand number of trips by month where customers were not happy.
- Also could give insight or a hidden pattern in dissatisfactory rides in a specific month or period which could be a number of factors like low rating because of AC was not on during summery time, cab reached late on pickup point due to traffic on a rainy day/season etc.
- Based on this analysis, instructions can be given to driver to make customers happy and take care of things which could lead to low customer rating.

OUTPUT:



cabrides

Tables

(3) + ↻

Filter...

bookings_detail

clickstream_data

datewise_total_bookings

1 SELECT

2 DATE_FORMAT(PICKUP_TIMESTAMP, 'yyyy-MM') TRIP_MONTH

3 , COUNT(BOOKING_ID) AS NO_OF_BOOKINGS

4 FROM

5 BOOKINGS_DETAIL

6 WHERE

7 RATING_BY_CUSTOMER < 2

8 GROUP BY

9 DATE_FORMAT(PICKUP_TIMESTAMP, 'yyyy-MM')

10 ORDER BY

11 TRIP_MONTH;

8.52s Database cabrides Type text ?

Query History

Saved Queries

Query Builder

Results (10)

	trip_month	no_of_bookings
1	2020-01	26
2	2020-02	16
3	2020-03	16
4	2020-04	21
5	2020-05	21
6	2020-06	14
7	2020-07	20
8	2020-08	32
9	2020-09	21
10	2020-10	15

VALIDATION: Exact Match

```
Total MapReduce CPU Time Spent: 7 seconds 970 msec
OK
2020-01 26
2020-02 16
2020-03 16
2020-04 21
2020-05 21
2020-06 14
2020-07 20
2020-08 32
2020-09 21
2020-10 15
```

- **Task 11:** Calculate the count of total iOS users.

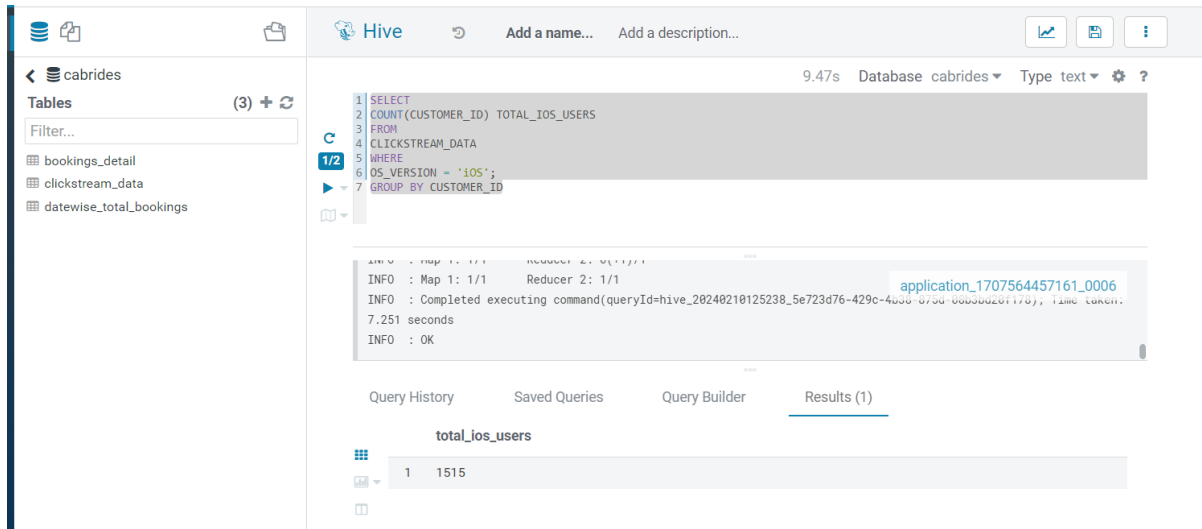
QUERY:

```
SELECT COUNT(CUSTOMER_ID) TOTAL_IOS_USERS
FROM CLICKSTREAM_DATA
WHERE OS_VERSION = 'iOS';
GROUP BY CUSTOMER_ID;
```

EXPLANATION:

- The DISTINCT option within the COUNT function is utilized to tally the distinct number of customers using iOS devices.
- The WHERE clause filters for events originating from iOS devices.
- The ORDER BY clause organizes the data in chronological order by the pickup month, using Trip month as an alias.
- This analysis offers a broad overview of the quantity or proportion of customers who use a particular type of device or operating system.
- For instance, should the company roll out updates to the iOS and Android mobile applications, this data could estimate the customer base that would be affected.

OUTPUT:



The screenshot shows the Hive query interface. On the left, a sidebar lists tables under the 'cabrides' database: 'bookings_detail', 'clickstream_data', and 'datewise_total_bookings'. The main area displays a SQL query:

```
1 SELECT
2 COUNT(DISTINCT CUSTOMER_ID) TOTAL_IOS_USERS
3 FROM
4 CLICKSTREAM_DATA
5 WHERE
6 OS_VERSION = 'iOS';
7 GROUP BY CUSTOMER_ID
```

Below the query, the execution progress is shown: 'Map 1: 1/1' and 'Reducer 2: 1/1'. The status is 'INFO : Completed executing command(queryId=hive_20240218125238_5e723d76-429c-4b3b-b73d-000300201176), time taken: 7.251 seconds'. The results are displayed in a table with one row:

total_ios_users
1515

VALIDATION: close to Match(1515) , since Kafka had extra 16 records compare to validation
You should get the count of all iOS users as 1503.